

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

MLOps Experimental Learning Assignment:

End-to-End ML Model Development, CI/CD, and Production Deployment Experimental Learning.

Objective:

Design, develop, and deploy a scalable, reproducible machine learning solution using modern MLOps best practices. The assignment emphasizes practical automation, experiment tracking, CI/CD pipelines, containerization, cloud deployment, and monitoring, mirroring real-world production scenarios.

Dataset: Title: Heart Disease UCI Dataset

- Source: UCI Machine Learning Repository
 - <https://archive.ics.uci.edu/ml/datasets/heart+Disease>
- CSV containing 14+ features (age, sex, blood pressure, cholesterol, etc.) and a binary target (presence/absence of heart disease).

Problem Statement:

Build a machine learning classifier to predict the risk of heart disease based on patient health data, and deploy the solution as a cloud-ready, monitored API.

Assignment Tasks:

1. Data Acquisition & Exploratory Data Analysis (EDA) [5 marks]

- Obtain the dataset (provide download script or instructions). Clean and preprocess the data (handle missing values, encode features). Perform EDA with professional visualizations (histograms, correlation heatmaps, class balance).

2. Feature Engineering & Model Development [8 marks]

- Prepare final ML features (scaling, encoding). Build and train at least two classification models (e.g., Logistic Regression and Random Forest). Document model selection and tuning process. Evaluate using cross validation and relevant metrics (accuracy, precision, recall, ROC-AUC).

3. Experiment Tracking [5 marks]

- Integrate MLflow (or similar) for experiment tracking. Log parameters, metrics, artifacts, and plots for all runs.

4. Model Packaging & Reproducibility [7 marks]

- Save the final model in a reusable format (e.g., MLflow, pickle, ONNX). Write a clean requirements.txt (or Conda env file). Provide preprocessing pipeline/transformers to ensure full reproducibility.

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

5. CI/CD Pipeline & Automated Testing [8 marks]

- Write unit tests for data processing and model code (Pytest or unit test). Create a GitHub Actions (or Jenkins) pipeline: Linting, unit test, and model training steps. Artifacts/logging for each workflow run.

6. Model Containerization [5 marks]

- Build a Docker container for the model serving API (Flask or FastAPI recommended). Expose /predict endpoint, accept JSON input, return prediction and confidence. Container must build and run locally with sample input.

7. Production Deployment [7 marks]

- Deploy the Dockerized API to a public cloud or local Kubernetes (GKE, EKS, AKS, or Minikube/Docker Desktop). Use a deployment manifest or Helm chart. Expose via Load Balancer or Ingress. Verify endpoints and provide deployment screenshots.

8. Monitoring & Logging [3 marks]

- Integrate logging of API requests. Demonstrate simple monitoring (Prometheus + Grafana or API metrics/logs dashboard).

9. Documentation & Reporting [2 marks]

- Submit a professional Markdown or PDF report including:
 - Setup/install instructions.
 - EDA and modelling choices.
 - Experiment tracking summary.
 - Architecture diagram.
 - CI/CD and deployment workflow screenshots.
 - Link to code repository.

Assignment Implementation:

Assignment Objective

This assignment requires the development of a complete MLOps workflow for a machine learning problem. The objective is not only to build a predictive model, but also to show how the solution can be prepared for repeatable training, automated testing, containerized deployment, monitoring, and documentation. The repository delivered here covers all of those aspects for a heart disease prediction system using the Cleveland heart disease dataset.

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN


Requirement 1: Data acquisition and exploratory data analysis

For this requirement focuses on obtaining the dataset, understanding its structure, and performing exploratory analysis before modelling.

Following things has been done,

- The raw dataset is stored in **data/raw/heart_cleveland.csv** and is loaded through the reusable preprocessing module in **src/data/preprocessing.py**
- The workflow includes a data quality report with missing value analysis, duplicate checks, feature statistics, and target distribution review.
- Missing values are handled systematically using median imputation for numeric features and mode imputation for categorical features.
- The repository also includes notebook-based exploratory work in notebooks/01_data_acquisition_eda.ipynb
- The solution shows how raw data is loaded, checked, cleaned, and inspected before any model is trained.
- The process is documented in code and supported by visual evidence.

Visual evidence:



```
heart_cleveland.csv X
data > raw > heart_cleveland.csv > data
1  age,sex,cp,trestbps,chol,fbs,restecg,thalach,exang,oldpeak,slope,ca,thal,target
2  63.0,1.0,1.0,145.0,233.0,1.0,2.0,150.0,0.0,2.3,3.0,0.0,6.0,0
3  67.0,1.0,4.0,160.0,286.0,0.0,2.0,108.0,1.0,1.5,2.0,3.0,3.0,1
4  67.0,1.0,4.0,120.0,229.0,0.0,2.0,129.0,1.0,2.6,2.0,2.0,7.0,1
5  37.0,1.0,3.0,130.0,250.0,0.0,0.0,187.0,0.0,3.5,3.0,0.0,3.0,0
6  41.0,0.0,2.0,130.0,204.0,0.0,2.0,172.0,0.0,1.4,1.0,0.0,3.0,0
7  56.0,1.0,2.0,120.0,236.0,0.0,0.0,178.0,0.0,0.8,1.0,0.0,3.0,0
8  62.0,0.0,4.0,140.0,268.0,0.0,2.0,160.0,0.0,3.6,3.0,2.0,3.0,1
9  57.0,0.0,4.0,120.0,354.0,0.0,0.0,163.0,1.0,0.6,1.0,0.0,3.0,0
10 63.0,1.0,4.0,130.0,254.0,0.0,2.0,147.0,0.0,1.4,2.0,1.0,7.0,1
11 53.0,1.0,4.0,140.0,203.0,1.0,2.0,155.0,1.0,3.1,3.0,0.0,7.0,1
12 57.0,1.0,4.0,140.0,192.0,0.0,0.0,148.0,0.0,0.4,2.0,0.0,6.0,0
13 56.0,0.0,2.0,140.0,294.0,0.0,2.0,153.0,0.0,1.3,2.0,0.0,3.0,0
14 56.0,1.0,3.0,130.0,256.0,1.0,2.0,142.0,1.0,0.6,2.0,1.0,6.0,1
15 44.0,1.0,2.0,120.0,263.0,0.0,0.0,173.0,0.0,0.0,1.0,0.0,7.0,0
16 52.0,1.0,3.0,172.0,199.0,1.0,0.0,162.0,0.0,0.5,1.0,0.0,7.0,0
17 57.0,1.0,3.0,150.0,168.0,0.0,0.0,174.0,0.0,1.6,1.0,0.0,3.0,0
18 48.0,1.0,2.0,110.0,229.0,0.0,0.0,168.0,0.0,1.0,3.0,0.0,7.0,1
```

RAW DATA


Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN



Name: RAJ KUMAR M
BITS ID: 2024AD05110
Course: MTECH AI & ML
Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
age														
sex	-0.10													
cp	0.10	0.01												
trestbps	0.28	-0.06	-0.04											
chol	0.21	-0.20	0.07	0.13										
fbs	0.12	0.05	-0.04	0.18	0.01									
restecg	0.15	0.02	0.07	0.15	0.17	0.07								
thalach	-0.39	-0.05	-0.33	-0.05	-0.00	-0.01	-0.08							
exang	0.09	0.15	0.38	0.06	0.06	0.03	0.08	-0.38						
oldpeak	0.20	0.10	0.20	0.19	0.05	0.01	0.11	-0.34	0.29					
slope	0.16	0.04	0.15	0.12	-0.00	0.06	0.13	-0.39	0.26	0.58				
ca	0.36	0.09	0.23	0.10	0.12	0.15	0.13	-0.26	0.15	0.30	0.11			
thal	0.13	0.38	0.27	0.13	0.01	0.07	0.02	-0.28	0.33	0.34	0.29	0.26		
target	0.22	0.28	0.41	0.15	0.09	0.03	0.17	-0.42	0.43	0.42	0.34	0.46	0.53	

- src/data/preprocessing.py
- data/raw/heart_cleveland.csv
- notebooks/01_data_acquisition_eda.ipynb
- screenshots/target_distribution.png
- screenshots/missing_values.png
- screenshots/correlation_heatmap.png

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110


Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Requirement 2: Feature engineering and model development

- The preprocessing pipeline defines numeric and categorical features and applies scaling and imputation in **src/data/preprocessing.py**
- The training workflow in **src/models/train.py** trains multiple classifiers, including logistic regression, random forest, gradient boosting, and support vector machine.
- Each model is evaluated using accuracy, precision, recall, F1-score, and ROC-AUC.
- The best-performing model is saved for later inference.
- The project demonstrates model experimentation rather than a single fixed classifier.
- The model development process is reusable and reproducible.

Visual evidence:



categorical_distributions


Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN



numeric_distributions



Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01


Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN


Pairplot of Key Features



Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01


Name: RAJ KUMAR M
BITS ID: 2024AD05110
Course: MTECH AI & ML
Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Logistic regression



screenshots > logistic_regression > ≡ classification_report.txt


		precision	recall	f1-score	support
1					
2					
3	0	0.93	0.82	0.87	33
4	1	0.81	0.93	0.87	28
5					
6	accuracy			0.87	61
7	macro avg	0.87	0.87	0.87	61
8	weighted avg	0.88	0.87	0.87	61
9					



Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01


Name: RAJ KUMAR M
BITS ID: 2024AD05110
Course: MTECH AI & ML
Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Random forest




classification_report.txt

screenshots > random_forest > classification_report.txt	
1	precision recall f1-score support
2	
3	0 0.93 0.85 0.89 33
4	1 0.84 0.93 0.88 28
5	
6	accuracy 0.89 61
7	macro avg 0.89 0.89 0.89 61
8	weighted avg 0.89 0.89 0.89 61




Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M
BITS ID: 2024AD05110
Course: MTECH AI & ML
Email: 2024AD05110@WILP.BITS-PILANI.AC.IN



Support Vector Machines (SVM)




Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M
BITS ID: 2024AD05110
Course: MTECH AI & ML
Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

≡ classification_report.txt X

screenshots > svm > ≡ classification_report.txt

		precision	recall	f1-score	support
1					
2					
3	0	0.87	0.79	0.83	33
4	1	0.77	0.86	0.81	28
5					
6	accuracy			0.82	61
7	macro avg	0.82	0.82	0.82	61
8	weighted avg	0.82	0.82	0.82	61



Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01


Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Gradient boosting:



```
≡ classification_report.txt X
```

```
screenshots > gradient_boosting > ≡ classification_report.txt
```

		precision	recall	f1-score	support
1					
2					
3	0	0.88	0.85	0.86	33
4	1	0.83	0.86	0.84	28
5					
6	accuracy			0.85	61
7	macro avg	0.85	0.85	0.85	61
8	weighted avg	0.85	0.85	0.85	61
9					



Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN



Relevant files in Repository:

- src/data/preprocessing.py
- src/models/train.py
- models/final_model.joblib
- screenshots/categorical_distributions.png
- screenshots/numeric_distributions.png
- screenshots/target_correlation.png

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

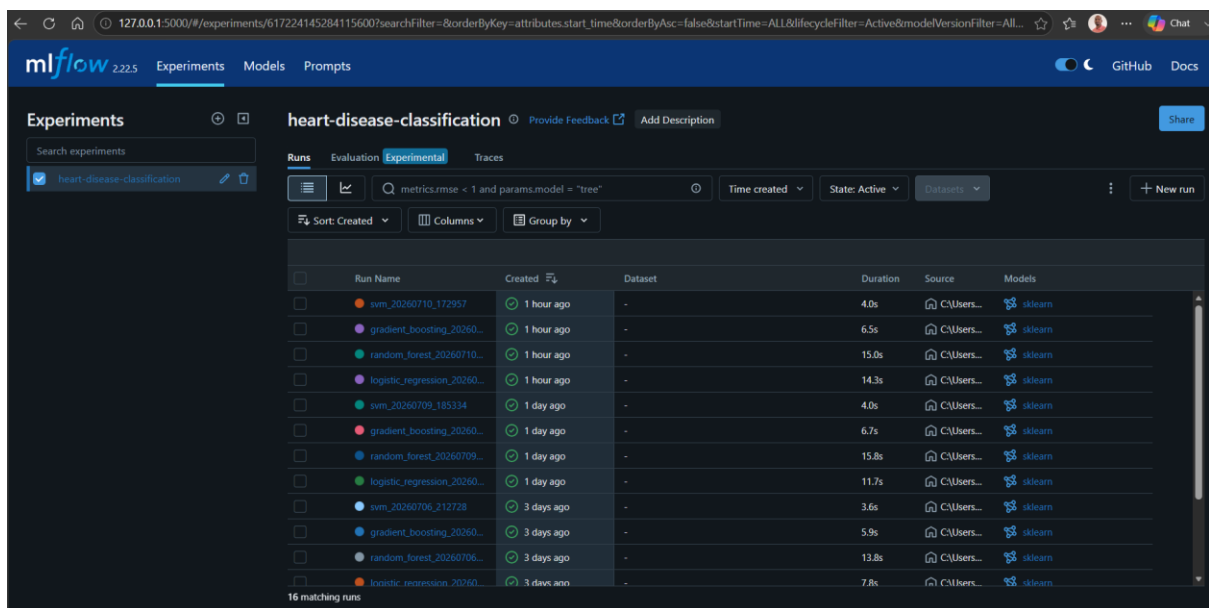
BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Requirement 3: Experiment tracking with MLflow

- The training script uses MLflow to log experiment metadata, model parameters, evaluation metrics, and generated plots.
- The local MLflow tracking server can be used to compare runs and inspect artifacts.
- Experiment artifacts are stored under mlruns
- Teams can reproduce and compare model runs over time.
- The reporting and performance comparison become traceable and auditable.



The screenshot shows the MLflow web interface for an experiment named 'heart-disease-classification'. The 'Runs' tab is active, displaying a table of 16 runs. The table columns are Run Name, Created, Dataset, Duration, Source, and Models. The runs are sorted by creation time, showing various models like svm, gradient_boosting, random_forest, and logistic_regression, each with its duration and source path.

Run Name	Created	Dataset	Duration	Source	Models
svm_20260710_172957	1 hour ago	-	4.0s	C:\Users...	sklearn
gradient_boosting_20260...	1 hour ago	-	6.5s	C:\Users...	sklearn
random_forest_20260710...	1 hour ago	-	15.0s	C:\Users...	sklearn
logistic_regression_20260...	1 hour ago	-	14.3s	C:\Users...	sklearn
svm_20260709_185334	1 day ago	-	4.0s	C:\Users...	sklearn
gradient_boosting_20260...	1 day ago	-	6.7s	C:\Users...	sklearn
random_forest_20260709...	1 day ago	-	15.8s	C:\Users...	sklearn
logistic_regression_20260...	1 day ago	-	11.7s	C:\Users...	sklearn
svm_20260706_212728	3 days ago	-	3.6s	C:\Users...	sklearn
gradient_boosting_20260...	3 days ago	-	5.9s	C:\Users...	sklearn
random_forest_20260706...	3 days ago	-	13.8s	C:\Users...	sklearn
logistic_regression_20260...	3 days ago	-	7.8s	C:\Users...	sklearn

Relevant files in Repository:

- src/models/train.py
- mlruns
- notebooks/02_model_training.ipynb

Requirement 4: Model packaging and reproducibility

- The trained model and fitted preprocessor are stored as joblib artifacts in **models/final_model.joblib** and **models/preprocessor.joblib**
- The code is organized so that data processing, model training, and prediction serving can be run independently.
- Dependency management is handled through **requirements.txt** and **setup.py**
- The trained assets can be reused for prediction without retraining.
- The environment is documented and can be recreated on another machine.

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Relevant files in Repository:

- models/final_model.joblib
- models/preprocessor.joblib
- requirements.txt
- setup.py

Requirement 5: CI/CD and automated testing

- Automated unit tests are included in **tests/test_api.py**, **tests/test_data_processing.py**, and **tests/test_model.py**.
- A GitHub Actions workflow is defined in **.github/workflows/ci-cd.yml** to run linting, tests, model training, and Docker build steps.
- The workflow is structured to provide reproducible validation each time the code changes.
- The repository includes automated checks that help catch regressions early.
- The workflow reflects a realistic CI/CD approach expected in MLOps assignments.

Verification evidence:

- If we ran the following command in the workspace:

pytest -q

Result:

```
PS C:\Users\Rajkumar Madhavan\OneDrive\Documents\GitHub\MLOPS_Assignment-1> (Set-ExecutionPolicy -Scope Process -ExecutionPolicy RemoteSigned) ; (& "c:\Users\Rajkumar Madhavan\OneDrive\Documents\GitHub\MLOPS_Assignment-1\.venv\Scripts\Activate.ps1")
(.venv) PS C:\Users\Rajkumar Madhavan\OneDrive\Documents\GitHub\MLOPS_Assignment-1> pytest -q
===== test session starts =====
platform win32 -- Python 3.11.9, pytest-9.1.1, pluggy-1.6.0
rootdir: C:\Users\Rajkumar Madhavan\OneDrive\Documents\GitHub\MLOPS_Assignment-1
configfile: pytest.ini
testpaths: tests
plugins: anyio-4.14.1, cov-7.1.0
collected 60 items

tests\test_api.py ..... [ 40%]
tests\test_data_processing.py ..... [ 71%]
tests\test_model.py ..... [100%]

===== warnings summary =====
..\..\Desktop\MLOPS\MLOPS-Assigment\venv\Lib\site-packages\fastapi\testclient.py:1
C:\Users\Rajkumar Madhavan\OneDrive\Desktop\MLOPS\MLOPS-Assigment\venv\Lib\site-packages\fastapi\testclient.py:1: StarletteDeprecationWarning: Using `httpx`
with `starlette.testclient` is deprecated; install `httpx2` instead.
  from starlette.testclient import TestClient as TestClient # noqa

-- Docs: https://docs.pytest.org/en/stable/how-to/capture-warnings.html
===== 60 passed, 1 warning in 17.41s =====
```

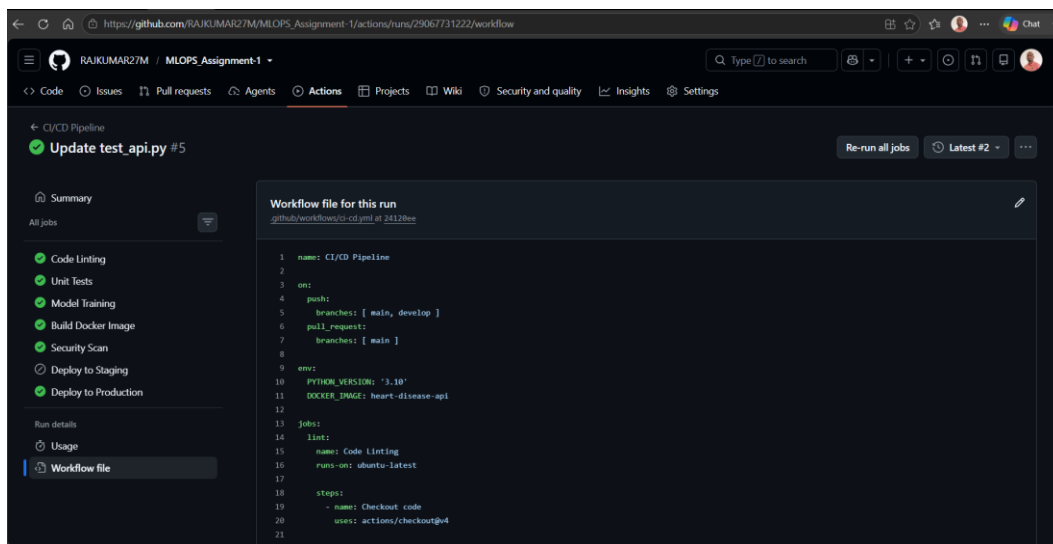
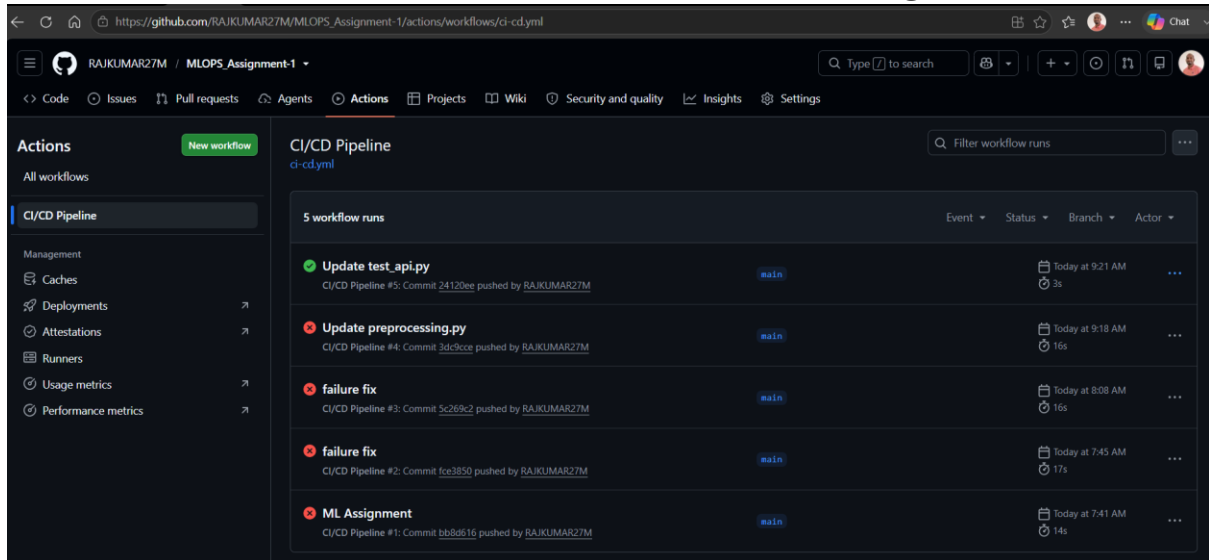

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN



Relevant files in Repository:

- .github/workflows/ci-cd.yml
- tests/test_api.py
- tests/test_data_processing.py
- tests/test_model.py

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Requirement 6: Model containerization

This requirement expects the application to be packaged in a portable container image.

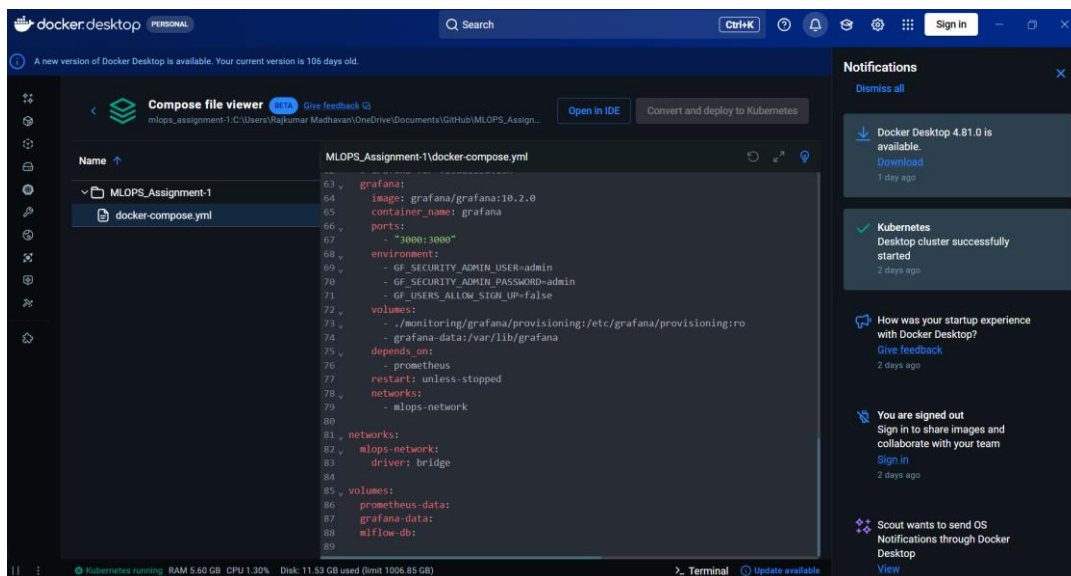
- The container definition is provided in **Dockerfile**.
- The Docker image installs dependencies, copies the application code, and starts the API service.
- Local orchestration and service composition are supported through **docker-compose.yml**
- The model API can be built and run as an isolated environment, which is essential for deployment consistency.

Relevant files in Repository:

- Dockerfile

- docker-co

mpose.yml



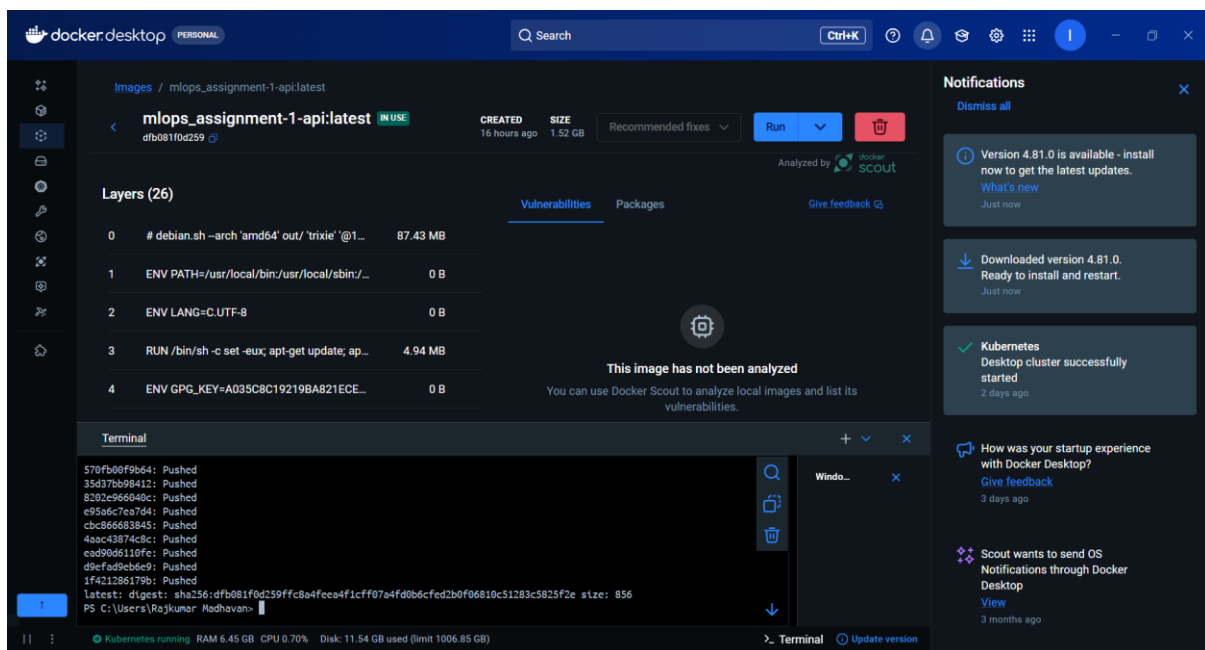
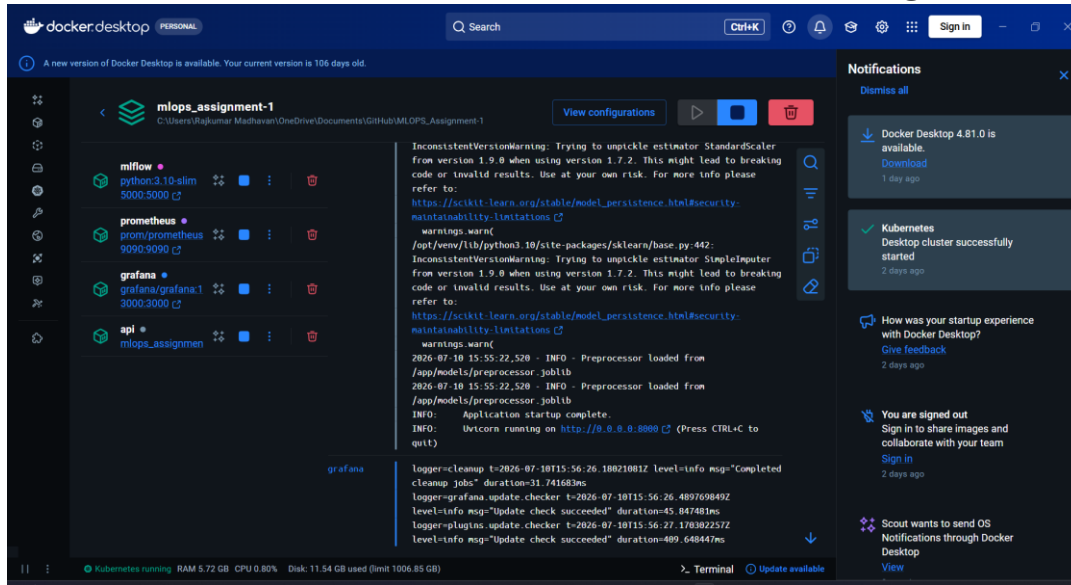
Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN



Requirement 7: Production-style deployment

- Kubernetes deployment manifests are included in **deployment/kubernetes/deployment.yaml**.
- A Helm chart is also provided in **deployment/helm/heart-disease-api**.
- The deployment uses a rolling update strategy, readiness and liveness probes, and resource requests and limits.
- The project demonstrates how the API could be served in a container orchestration environment, not only as a local script.

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Relevant files in Repository:

- deployment/kubernetes/deployment.yaml
- deployment/kubernetes/service.yaml
- deployment/kubernetes/ingress.yaml
- deployment/helm/heart-disease-api

Requirement 8: Monitoring and logging

- The API in **src/api/main.py** includes request logging and prediction endpoints.
- Prometheus metrics are exposed through the application and configured in **monitoring/prometheus/prometheus.yml**.
- Grafana provisioning and dashboard configuration are available in **monitoring/grafana**.
- The application is prepared for production observability and can be monitored for usage and performance.

Relevant files in Repository:

- src/api/main.py
- monitoring/prometheus/prometheus.yml
- monitoring/grafana

Requirement 9: Documentation and reporting

- The main project documentation is available in **README.md**
- This detailed report is provided in **FINAL_REPORT.md**
- A shorter submission summary is available in **GITHUB_SUBMISSION_SUMMARY.md**
- The repository also preserves screenshots and artifacts in **screenshots**

Code Repository Link

GitHub repository: https://github.com/RAJKUMAR27M/MLOPS_Assignment-1

Link to Video: <https://drive.google.com/file/d/1IjVZdUONCeYiAnmgSpUvEuxlNhC50K-R/view?usp=sharing>

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M


BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Project Architecture Summary

The repository is organized in a way that reflects a realistic MLOps pipeline:



- Data layer: **src/data**
- Model training layer: **src/models**
- Serving layer: **src/api**
- Testing layer: **tests**
- Deployment and orchestration: **deployment**
- Monitoring: **monitoring**

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M


BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Code Repository Link

GitHub repository: https://github.com/RAJKUMAR27M/MLOPS_Assignment-1



How to Run the Project

- **Prerequisites**

- Install Python 3.11+
- Install Git and VS Code
- Ensure We are in the project root folder

- **Create and activate the environment**

Open PowerShell in the project folder and run below commands one by one:

- `python -m venv .venv`
- `.\.venv\Scripts\Activate.ps1`
- `python -m pip install --upgrade pip`
- `pip install -r requirements.txt`

- **Explore the dataset and understand the problem**

Review the raw data file at **data/raw/heart_cleveland.csv**

Open the notebooks in **notebooks** for EDA and preprocessing understanding

Inspect the preprocessing logic in **src/data/preprocessing.py**

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

- **Run preprocessing and prepare the data**

Run below command:

```
python -m src.data.preprocessing
```

This step loads the data, handles missing values, encodes features, and prepares training data.

- **Train the machine learning models**

Run below commands:

```
python -m src.models.train
```

This trains multiple classifiers, evaluates them, and saves the best model plus preprocessor to **models**.

- **Track experiments with MLflow**

Start the MLflow UI:

```
python -m mlflow ui --host 127.0.0.1 --port 5000
```

Then open:

<http://127.0.0.1:5000>

We can compare training runs and inspect metrics/artifacts stored in **mlruns**.

- **Run the automated tests**

Run below command:

```
pytest -q
```

This verifies preprocessing, training, and API behaviour.

- **Start the prediction API locally**

Run below command:

```
python -m uvicorn src.api.main:app --host 127.0.0.1 --port 8000
```

Then open:

<http://127.0.0.1:8000/docs>

We can test the prediction endpoint from the Swagger UI.

- **Containerize the application**

Build the Docker image:

Run below command in powershell but make sure docker running

```
docker build -t heart-disease-api .
```

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

- **Run the container:**

`docker run -p 8000:8000 heart-disease-api`

- **Run with Docker Compose**

Use:

`docker-compose up`

This starts the services defined in **docker-compose.yml**

- **Deploy with Kubernetes or Helm**

Deployment assets are available in **deployment/kubernetes** and **deployment/helm/heart-disease-api**

Example:

In powershell


`kubectl apply -f deployment/kubernetes/`

- **Monitor the application**

Monitoring configuration is available in **monitoring/prometheus** and **monitoring/grafana**

- Prometheus scrapes the metrics endpoint
- Grafana visualizes the service metrics

Swagger UI Screenshots:



Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Root

GET / Root

Root endpoint with API Information.

Parameters

No parameters

Try it out

Responses

Curl

```
curl -X 'GET' \
  'http://127.0.0.1:8000/' \
  -H 'accept: application/json'
```

Request URL

http://127.0.0.1:8000/

Server response

Code	Details
200	Response body

127.0.0.1:8000/docs#/Root

Code Details

200

Response body

```
{
  "name": "Heart Disease Prediction API",
  "version": "1.0.0",
  "description": "ML-powered heart disease risk prediction",
  "endpoints": {
    "health": "/health",
    "predict": "/predict",
    "batch_predict": "/predict/batch",
    "model_info": "/model/info",
    "metrics": "/metrics",
    "docs": "/docs"
  }
}
```

Response headers

```
content-length: 264
content-type: application/json
date: Fri, 10 Jul 2026 12:06:42 GMT
server: uvicorn
```

Responses

Code	Description	Links
200	Successful Response	No links

Media type

application/json

Controls Accept header.

Example Value | Schema

```
"string"
```

127.0.0.1:8000/docs#/Health

Health

GET /health Health Check

Health check endpoint.

Returns the current health status of the API.

Parameters

No parameters

Try it out

Responses

Code	Description	Links
200	Successful Response	No links

Media type

application/json

Controls Accept header.

Example Value | Schema

```
{
  "model_loaded": true,
  "model_version": "1.0.0",
  "status": "healthy"
}
```

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Model

GET

/model/info

Get Model Info

Get model information.
Returns details about the loaded model.

Parameters

Try it out

No parameters

Responses

Curl

```
curl -X 'GET' \
  'http://127.0.0.1:8000/model/info' \
  -H 'accept: application/json'
```

Request URL

http://127.0.0.1:8000/model/info

Request URL

http://127.0.0.1:8000/model/info

Server response

Code

Details

200

Response body

```
{
  "model_name": "random_forest",
  "model_version": "1.0.0",
  "features": [
    "age",
    "sex",
    "cp",
    "trestbps",
    "chol",
    "fbs",
    "restecg",
    "thalach",
    "exang",
    "oldpeak",
    "slope",
    "ca",
    "thal"
  ],
  "created_at": "2026-07-10T17:30:01.177718"
}
```

Download

Response headers

```
content-length: 288
content-type: application/json
date: Fri, 10 Jul 2026 12:07:15 GMT
server: uvicorn
```

Responses

Code

Description

Links

200

Successful Response

No links

Media type

application/json

Controls Accept header.

Example Value | Schema

```
{
  "created_at": "2024-01-15T10:30:00",
  "features": [
    "age",
    "sex",
    "cp",
    "trestbps",
    "chol",
    "fbs",
    "restecg",
    "thalach",
    "exang",
    "oldpeak",
    "slope",
    "ca",
    "thal"
  ],
  "model_name": "random_forest",
  "model_version": "1.0.0"
}
```

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Prediction

POST

/predict

Predict

Predict heart disease risk for a single patient.

- age: Age in years (20-100)
- sex: Sex (1 = male, 0 = female)
- cp: Chest pain type (0-3)
- trestbps: Resting blood pressure (mm Hg)
- chol: Serum cholesterol (mg/dl)
- fbis: Fasting blood sugar > 120 mg/dl (1 = true, 0 = false)
- restecg: Resting ECG results (0-2)
- thalach: Maximum heart rate achieved
- exang: Exercise induced angina (1 = yes, 0 = no)
- oldpeak: ST depression induced by exercise
- slope: Slope of peak exercise ST segment (0-2)
- ca: Number of major vessels (0-4)
- thal: Thalassemia (0-3)

Parameters

No parameters

Try it out

Request body

required

application/json

Example Value

Schema

```
{
  "age": 63,
  "ca": 0,
  "chol": 233,
  "cp": 3,
  "exang": 0,
  "fbis": 1,
  "oldpeak": 2.3,
  "restecg": 0,
  "sex": 0,
  "slope": 0,
  "thal": 1,
  "thalach": 150,
  "trestbps": 145
}
```

Responses

Code	Description	Links
200	Successful Response	No links
	Media type application/json Controls Accept header. Example Value Schema <pre>{ "confidence": 0.85, "model_version": "1.0.0", "prediction": 1, "probability": 0.85, "risk_level": "High" }</pre>	
422	Validation Error	No links
	Media type application/json Example Value Schema <pre>{ "detail": [{ "loc": ["string", 0], "msg": "string", "type": "string", "input": "string", "ctx": {} }] }</pre>	

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Responses		
Code	Description	Links
200	Successful Response	No links
Media type application/json Controls Accept header. Example Value Schema <pre>{ "predictions": [{ "confidence": 0.85, "model version": "1.0.0", "prediction": 1, "probability": 0.85, "risk_level": "High" }], "count": 0}</pre>		
422	Validation Error	No links
Media type application/json Example Value Schema <pre>{ "detail": [{ "loc": ["string", 0], "msg": "string", "type": "string", "input": "string", "ctx": {} }]}</pre>		

Monitoring

GET /metrics Metrics	
Prometheus metrics endpoint. Returns metrics in Prometheus format for scraping.	
Parameters	Try it out
No parameters	
Responses	
Curl <pre>curl -X 'GET' \ 'http://127.0.0.1:8000/metrics' \ -H 'accept: application/json'</pre> Request URL http://127.0.0.1:8000/metrics	

Machine Learning Operations (MLOps) AIMLCZG523 Assignment 01

Name: RAJ KUMAR M

BITS ID: 2024AD05110

Course: MTECH AI & ML

Email: 2024AD05110@WILP.BITS-PILANI.AC.IN

Request URL
http://127.0.0.1:8000/metrics

Server response

Code	Details
200	Response body <pre># HELP python_gc_objects_collected_total Objects collected during gc # TYPE python_gc_objects_collected_total counter python_gc_objects_collected_total{generation="0"} 978.0 python_gc_objects_collected_total{generation="1"} 719.0 python_gc_objects_collected_total{generation="2"} 25.0 # HELP python_gc_objects_uncollectable_total Uncollectable objects found during GC # TYPE python_gc_objects_uncollectable_total counter python_gc_objects_uncollectable_total{generation="0"} 0.0 python_gc_objects_uncollectable_total{generation="1"} 0.0 python_gc_objects_uncollectable_total{generation="2"} 0.0 # HELP python_gc_collections_total Number of times this generation was collected # TYPE python_gc_collections_total counter python_gc_collections_total{generation="0"} 434.0 python_gc_collections_total{generation="1"} 39.0 python_gc_collections_total{generation="2"} 3.0 # HELP python_info Python platform information # TYPE python_info gauge python_info{implementation="CPython",major="3",minor="11",patchlevel="9",version="3.11.9"} 1.0 # HELP prediction_requests_total Total number of prediction requests # TYPE prediction_requests_total counter prediction_requests_total{endpoint="predict",status="success"} 1.0 prediction_requests_total{endpoint="batch_predict",status="success"} 1.0 # HELP prediction_requests_created Total number of prediction requests # TYPE prediction_requests_created gauge prediction_requests_created{endpoint="predict",status="success"} 1.7836852715547228e+09 prediction_requests_created{endpoint="batch_predict",status="success"} 1.783685352716777e+09 # HELP prediction_request_latency_seconds Request latency in seconds # TYPE prediction_request_latency_seconds histogram prediction_request_latency_seconds_bucket{endpoint="predict",le="0.005"} 0.0 prediction_request_latency_seconds_bucket{endpoint="predict",le="0.01"} 0.0</pre> Response headers <pre>content-length: 5049 content-type: text/plain; version=1.0.0; charset=utf-8 date: Fri, 10 Jul 2026 12:17:07 GMT server: uvicorn</pre>
200	Successful Response

Responses

Code	Description	Links
200	Successful Response	No links

API deployed Publicly:

We can access it through public API URL <https://mlops-assignment-1-api-latest.onrender.com/>

Interactive API Docs: <https://mlops-assignment-1-api-latest.onrender.com/docs>

Conclusion

This project successfully demonstrates an end-to-end MLOps solution for heart disease prediction. It includes data preparation, model training, experiment tracking, testing, packaging, containerization, deployment assets, monitoring, and documentation. The repository is structured in a way that is suitable for both academic submission and practical industrial-style deployment demonstration.

Code Repository Link

GitHub repository: https://github.com/RAJKUMAR27M/MLOPS_Assignment-1

public API URL <https://mlops-assignment-1-api-latest.onrender.com/>

Interactive API Docs: <https://mlops-assignment-1-api-latest.onrender.com/docs>

Link to Video: <https://drive.google.com/file/d/1IjVZdUONCeYiAnmgSpUvEuxlNHc50K-R/view?usp=sharing>